



PEP003 Advanced Dependencies Update

PEP003 – Advanced Dependencies

- ▶ Full dependency graph, instead of single list of dependencies
- ▶ Reuse existing installations
- ▶ Any compatible bundle may satisfy a dependency
- ▶ Pass a dependency output to another dependency
- ▶ Only installs bundles. Lifecycle management is a future PEP

getporter.org/proposals/src/pep/003-dependency-namespaces-and-labels.md


Current Status

- ▶ Planning phase only as v1 of Porter and the Operator is our #1 priority
- ▶ We have an approved design for the user experience
 - ▶ Authoring a bundle with dependencies
 - ▶ Installing a bundle with dependencies
- ▶ Initial POC work is underway and incomplete
- ▶ Implementation is complex and is critical for our roadmap
- ▶ PEP needs implementation questions and decisions documented

The Road So Far...

- ▶ Schema changes are already in v1 release
- ▶ Goal: incremental development, no big bangs, shipped in v1.x
- ▶ Refactoring in-progress
 - ▶ ☒ Query installations (prep to match existing installations)
 - ▶ ☒ Passed context and config to the porter.yaml to bundle.json converter
 - ▶ ☒ List tags from an OCI repository (prep to resolve latest available tag)
 - ▶ ☐ Add feature flag
 - ▶ ☐ Move existing dependencies into v1 package
- ▶ Playground branch where I'm vetting implementation decisions

Playground Branch

- ▶ Define advanced dependencies in porter.yaml
- ▶ Represent as a new CNAB extension in bundle.json
- ▶ Generate a bundle dependencies graph
- ▶ Solve the graph
 - ▶ Reuse existing installations
 - ▶ Pick latest matching bundle from a registry
 - ▶ Check if a bundle satisfies a bundle interface
- ▶  Generate an execution plan
 - ▶ Execution Order
 - ▶ Output passing
 - ▶ Identify what can run in parallel
 - ▶ JIT secret injection

Workflows

- ▶ Execution plans are converted to a **Workflow**
 - ▶ Defines the tasks, inputs, and execution order
 - ▶ Decides what it will run in series or parallel
- ▶ Workflow is how that execution plan will run on a particular engine
- ▶ Stateful and persisted to Porter's database
- ▶ Users can see what has run, failed, is running, is pending
- ▶ Workflows can be restarted/retried when it fails

Let's talk workflow engines

- ▶ Executing dependencies (and bundles themselves!) are workflows
- ▶ Evolution of our internal workflow implementation
- ▶ Future work to secure mixins will use workflows too

Can we use third-party workflow engines?

- ▶ Existing workflow engines
 - ▶ Kubernetes: Argo Workflow, Brigade
 - ▶ Docker: Cadence Workflow(?)
- ▶ More research needed
 - ▶ How to use JIT secret injection with third-party workflow engines
 - ▶ Do not copy and persist secrets out of a secret store
 - ▶ Understand security implications
 - ▶ Verify that our workflow plugin interface is possible with existing engines

What's next?

- ▶ Refactoring and prep work continues
- ▶ Update PEP with implementation concepts and open decisions
- ▶ **Need your help vetting workflow engines and security**